



for Platform Builders

CoreOS Fest Berlin  
May 9-10 2016

Luke Bond  
@lukeb0nd

## # IN THIS TALK

- Some history
- Some opinions about rkt and Docker in relation to building platforms with them
- A practical / demo section on how to use rkt

This is essentially "rkt 101" with a few of my opinions thrown in.

## # WHO AM I?

- I'm a backend developer, DevOps-curious
- Mostly I do Node.js and Docker
- Built some things with Docker on CoreOS
  - Including a thing called "Paz" - [<http://paz.sh>]

I work for YLD.io, a London-based software engineering consultancy that specialises in Node.js, Docker and React.

## # HISTORY - ORIGINS OF RKT

- CoreOS felt Docker was no longer "the simple composable building block [they] had envisioned"
- rkt announced alongside the appc set of specifications
  - Specifications for runtime, image and discovery
- Focus was on the composability, security, image distribution, and openness of containers
- Docker's daemon and monolithic CLI tool make composability a problem
  - Leading to wrapping the Docker CLI by some tool makers (pre plugins)

## # HISTORY - STANDARDS ETC.

- Open Container Initiative (initially the Open Container Project) launched on June 22nd, 2015
- Docker's container image format and runtime donated to the project
- Appc shook things up enough to get Docker to commit to open standards
- Progress on standards appeared slow initially, and limited to the runtime
  - Recently work on the image and distribution has begun
  - See @vbatts' talk later today for an update

For a more detailed coverage of this history, check out this talk:  
<https://skillsmatter.com/skillscasts/7443-introduction-to-rkt-luke-bond>

## # DOCKER & RKT

- rkt doesn't have a daemon architecture like Docker does
- Docker's daemon architecture is a key differentiator for Docker
- This difference is what makes rkt great for building platforms

## # THE DOCKER DAEMON

- The Docker daemon model makes for a great integrated solution if you don't mind lock-in
- Docker's lack of composability a limiting factor for platform-builders
  - I know this from building Paz
- Those building platforms on Docker's tools are in a conundrum
  - Go all in on Compose & Swarm and have no differentiating tech
  - Build only on the daemon (i.e. build your own multi-host networking), and suffer lack of control over processes
- This is also about who owns PID1- systemd, Docker, ... unikernels?

## # RKT FOR PLATFORM BUILDERS

This is what makes `rkt` so attractive to platform builders.

```
> _There are really two buckets of users for Rocket and they
> could both be considered "platform builders."_
>
> The first set of platform builders are companies like Cloud
> Foundry, Mesosphere, or cloud service providers (Amazon Web
> Services, Google, Rackspace) that are building a platform as
> their product. Rocket allows them to add containers to their
> platform while keeping the rest of what they do today.
>
> The second set are enterprises that already have an existing
> environment and want to add containers to it. These would
> typically be large enterprises that have already invested in
> their own internal platform and want to layer in containers.
-> -- Alex Polvi*, CoreOS, 7/1/2015 on readwrite.com <-
```



## # HISTORY - THE CONTAINER WARS

- Docker is *\_making a platform play\_* \*

> My understanding is that *\_the Docker Platform will be a choice  
> for companies that want vSphere for containers\_* - that is, they  
> want a whole platform off the shelf.

-> -- *\*Alex Polvi\**, CoreOS, 7/1/2015 on readwrite.com <-

\* <https://blog.docker.com/2014/09/docker-closes-40m-series-c-led-by-sequoia/>

## # RKT FOR PLATFORM-BUILDING

- Tools should follow the Unix philosophy
- The Docker monolith does too much; as root
- rkt is composable
- Much of Docker's functionality is done well by other tools
  - Linux init systems like systemd manage processes perfectly well
  - tar/gzip are fine for packaging and distribution
  - etc.

## # RKT IN PRACTICE

- `rkt` is a container runtime, not an image build tool
- `rkt` started out as an appc reference implementation
- Therefore it natively runs appc images:

> `acbuild` is a command line utility to build and modify App  
> Container Images (ACIs), the container image format defined in  
> the App Container (appc) spec.

`acbuild` commands are analogous to lines of a `Dockerfile`.

A shell script of `acbuild` commands is the appc equivalent of a `Dockerfile`.

Let's try an example.

```
# RKT IN PRACTICE
```

```
## Building Images with acbuild
```

Let's take the most basic of example applications:

```
$ cat hello.c
#include <stdio.h>
int main (int argc, char** argv) {
    printf("Hello, world!\n");
    return 0;
}
$ gcc -o hello -static hello.c
```

## # RKT IN PRACTICE

We could use the following script to make an appc image:

```
$ cat appc-hello.sh
#!/usr/bin/env bash
acbuild begin
acbuild set-name hello
acbuild copy hello /app/hello
acbuild set-working-directory /app
acbuild set-exec -- /app/hello
acbuild write --overwrite hello-latest-linux-amd64.aci
```

Run the build script to build the image:

```
$ ./appc-hello.sh
$ ls -l hello-latest-linux-amd64.aci
-rw-r--r-- 1 luke users 331282 May  6 01:08 hello-latest-linux-amd64.aci
```

## # RKT IN PRACTICE

### ## Launching Containers

Let's launch the container:

```
$ sudo rkt --insecure-options=image run hello-latest-linux-amd64.aci
image: using image from local store for image name coreos.com/rkt/stage1-coreos:1.4.0
image: using image from file hello-latest-linux-amd64.aci
networking: loading networks from /etc/rkt/net.d
networking: loading network default with type ptp
[28959.906955] hello[4]: Hello, world!
```

As in Docker, the container is still present after execution has completed:

```
$ sudo rkt list
UUID          APP      IMAGE NAME  STATE    CREATED          STARTED          NETWORKS
3fa4bc21     hello   hello      exited  3 minutes ago   3 minutes ago
```

```
# RKT IN PRACTICE
```

```
## Housekeeping
```

Cleaning up completed containers is simple:

```
$ sudo rkt gc
```

Similarly for images that are no longer needed:

```
$ sudo rkt image gc
```

```
# RKT IN PRACTICE
```

```
## More Examples
```

Let's run something with networking:

```
$ git clone https://github.com/lukebond/demo-api.git
$ cd demo-api
$ sudo ./appc.sh
$ sudo rkt --insecure-options=image run demo-api-latest-linux-amd64.aci
image: using image from local store for image name coreos.com/rkt/stage1-coreos:1.4.0
image: using image from file demo-api-latest-linux-amd64.aci
image: searching for app image quay.io/coreos/alpine-sh
image: remote fetching from URL "https://quay.io/c1/aci/quay.io/coreos/alpine-sh/latest/aci"
Downloading ACI: [=====] 2.65 MB/2.65 MB
networking: loading networks from /etc/rkt/net.d
networking: loading network default with type ptp
```



## # RKT IN PRACTICE

In another terminal, check its status and make a request to it:

```
$ sudo rkt list
```

UUID	APP	IMAGE NAME	STATE	CREATED	STARTED	NETWORKS
3fa4bc21	hello	hello	exited	17 minutes ago	17 minutes ago	
5f9095eb	demo-api	lukebond/demo-api	running	52 seconds ago	52 seconds ago	default:ip4=1

```
$ curl 172.16.28.9:9000
```

```
"Hello, world 172.16.28.9!"
```

Success!

```
# RKT IN PRACTICE
```

```
## Logs
```

To access the logs for our containers, we use systemd's journalctl, like so:

```
$ machinectl list
```

MACHINE	CLASS	SERVICE
rkt-5f9095eb-5ad0-4cb1-b54f-729b1a3b3217	container	nspawn

```
1 machines listed.
```

```
$ sudo journalctl -M rkt-5f9095eb-5ad0-4cb1-b54f-729b1a3b3217
```

```
-- Logs begin at Fri 2016-05-06 01:37:58 BST, end at Fri 2016-05-06 01:37:5
```

```
...
```

```
# RKT IN PRACTICE
```

## ``` ## Stopping Containers ```

Killing containers is a bit of a hassle at the moment:

```
$ machinectl kill rkt-5f9095eb-5ad0-4cb1-b54f-729b1a3b3217
```

This is because the implementation we're using leverages systemd-nspawn to launch containers.

A native rkt command for stopping containers is reportedly coming in a future release.

## # RKT IN PRACTICE

### ## Signing images

- The `--insecure-options=image`` argument to ``rkt run`` disables signature verification
- But signing is important to ensure the artefact hasn't been tampered with
- Signing images is easily done using standard gpg tools.
- Instructions can be found here:

<https://coreos.com/rkt/docs/0.5.4/signing-and-verification-guide.html>

# RKT IN PRACTICE

## ## Converting Docker Images

The appc tool ``docker2aci`` can be used to download Docker images and convert them to appc's ACI format. Get the tool [[here](https://github.com/appc/docker2aci)](https://github.com/appc/docker2aci).

Converting a Docker image is as simple as:

```
$ docker2aci docker://lukebond/demo-api
```

It will squash the Docker layers into one ACI file.

To keep them separate, pass ``--nosquash``; it will set the correct dependencies between the layers.

```
# RKT IN PRACTICE
```

```
## Running Docker Images Directly
```

rkt can also run Docker images directly:

```
$ sudo rkt run --insecure-options=image docker://lukebond/demo-api
```

The "insecure" option is required here because Docker doesn't support the same image signature verification that rkt does.

Under the hood this is simply using ``docker2aci`` to first convert to ACI.

## # RKT IN PRACTICE

### ## Image Discovery and Distribution

- Image discovery in rkt follows the appc spec
- Whereas Docker uses a registry (defaulting to the Docker Hub), rkt uses a combination of HTTPS and HTML meta tags via a discovery URL
- This is best illustrated by an example (for CoreOS' Etcd):

```
$ curl -sL https://quay.io/coreos/etcd | grep meta | grep discovery
<meta name="ac-discovery" content="quay.io https://quay.io/c1/aci/{name}/{version}/{ext}">
<meta name="ac-discovery-pubkeys" content="quay.io https://quay.io/aci-signing-key">
```

- The content attributes are templated locators

For a full run-down of how appc image discovery works, read the specification

<https://github.com/appc/spec/blob/master/spec/discovery.md>

## # PODS

- The term "pod" is was popularized by the Kubernetes project
- A pod is a collection of applications should be scheduled as a unit
- The appc spec defines a pod as:
  - > "the deployable, executable unit...a list of apps that will be launched
  - > together inside a shared execution context, which includes network
  - > configuration and isolators. Whether you're running one process or multiple,
  - > rkt still considers it a pod.
- Pods are 1st-class citizens in rkt
- So far we've been talking about running containers; really we were running pods



## # PODS

Let's try a networked example; a trivial extension of demo-api with a Redis counter.

```
$ git clone https://github.com/lukebond/demo-api-redis.git
$ cd demo-api-redis
$ sudo ./appc.sh
$ sudo rkt run --volume volume--var-lib-redis,kind=host,source=/var/lib/redis quay.io/quay
  --insecure-options=image --port=http:9000 --set-env REDIS_HOST=localhost \
  demo-api-redis-latest-linux-amd64.aci
```

This will launch one pod, which contains two processes- Redis and a demo app, mapping necessary ports and volumes.

## # PODS

```
$ sudo rkt list
```

UUID	APP	IMAGE NAME	STATE	CREATED	STARTED	NE
e16bafd0	redis	quay.io/quay/redis:latest	running	6 seconds ago	6 seconds ago	de
	demo-api-redis	lukebond/demo-api-redis				

```
$ curl 172.16.28.6:9000
```

```
"Hello, world 172.16.28.6! 1 hits."
```

```
$ curl 172.16.28.6:9000
```

```
"Hello, world 172.16.28.6! 2 hits."
```

```
$ curl 172.16.28.6:9000
```

```
"Hello, world 172.16.28.6! 3 hits."
```

Success!

## # CONCLUSION

- I hope that has brought you up to speed with the basics of rkt
- If you're building a container-based platform, I recommend using rkt
- Docker is great but it will make your life harder for this use case

## # FURTHER READING

- ## Latest rkt features, including "fly"  
<https://coreos.com/blog/rkt-0.15.0-introduces-rkt-fly.html>
- ## rkt announcement - CoreOS blog  
<https://coreos.com/blog/rocket/>
- ## Open Container Project announcement - Docker blog  
<https://blog.docker.com/2015/06/open-container-project-foundation/>
- ## Weaveworks' Analysis of the OCP announcement - Weaveworks blog  
<http://www.weave.works/docker-open-container-project-please-make-it-awesome/>
- ## CoreOS' co-announcement of the OCP - CoreOS blog  
<https://coreos.com/blog/app-container-and-the-open-container-project/>
- ## OCP Progress report including name change to OCI - Docker blog  
<http://blog.docker.com/2015/07/open-container-format-progress-report/>
- ## CoreOS on OCI, appc and other standards - CoreOS blog  
<https://coreos.com/blog/making-sense-of-standards/>
- ## My rkt article on the Codeship blog  
<https://blog.codeship.com/getting-started-rkt/>
- ## OCI Image Spec project announcement  
<https://www.opencontainers.org/news/blogs/2016/04/new-image-specification-project-container>
- ## OCI Image Spec project on the CoreOS blog  
<https://coreos.com/blog/oci-image-specification.html>
- ## My talk at Docker London re rkt  
<https://skillsmatter.com/skillscasts/7443-introduction-to-rkt-luke-bond>